

Project 2: User Programs

Presented by

Xiaomo Liu

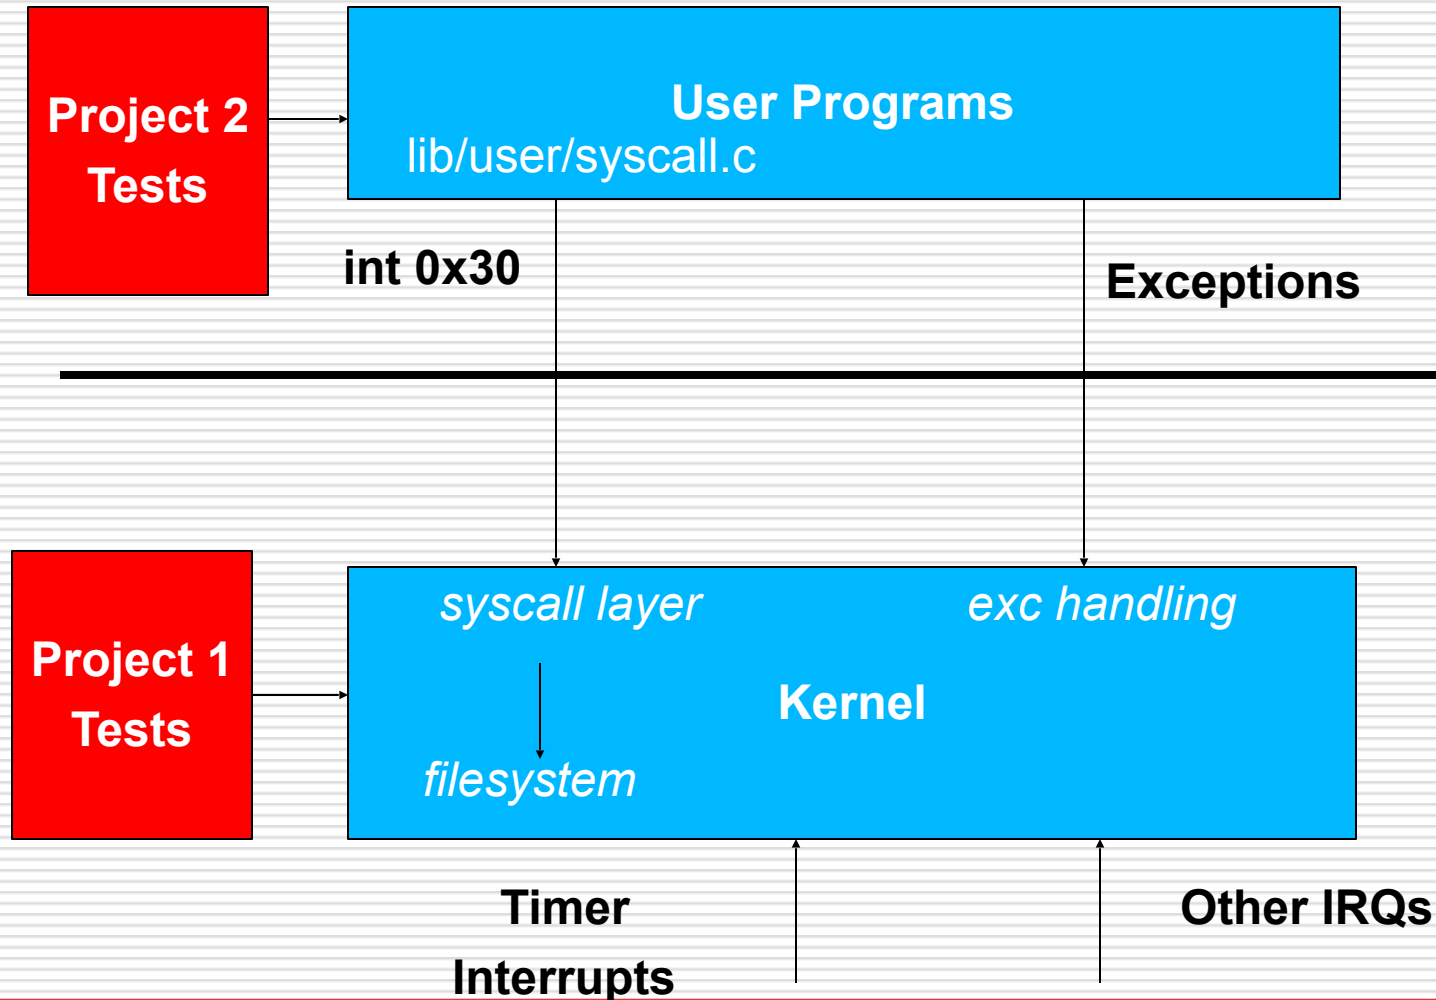
1 Oct 2009

(update Min Li's slides)

Till now ...

- All code part of Pintos Kernel
- Code compiled directly with the kernel
 - This required that the tests call some functions whose interface should remain unmodified
- From now on, run user programs on top of kernel
 - Freedom to modify the kernel to make the user programs work

Why Project 2 is not Project 1?



Sample User Program in C

- In C, a user program test.c can pass argument
 - ```
int main(int argc, char* argv[])
{
 for(int i=0; i<argc; i++)
 {
 char* arg = argv[i];
 }
}
```
  - `./test arg1 arg2 ...`

# Sample User Program in C

---

- test.c can call system libraries

- #include <stdio.h>

- int main()

- {

- FILE\* p\_file = fopen("myfile.txt","w");

- if (p\_file != NULL) fputs("fopen", p\_file);

- fclose(p\_file);

- }

- Get fopen, fputs, fclose by system calls

- Pintos need you to implement

- Argument passing

- System calls

# Using the File system

---

- ❑ May need to interact with file system
- ❑ Do not modify the file system!
  
- ❑ Certain limitations (till Project 4)
  - No internal synchronization
  - Fixed file size
  - No subdirectories
  - File names limited to 14 chars
  - System crash might corrupt the file system
  
- ❑ Files to take a look at: *'filesys.h'* & *'file.h'*

# Some commands

---

- Creating a simulated disk
  - `pintos-mkdisk filesys.dsk --fileysys-size=2`
- Formatting the disk
  - `pintos -f -q`
  - This will only work after your kernel is built !
- Copying the program into the disk
  - `pintos -p ../../examples/echo -a echo -- -q`
- Running the program
  - `pintos -q run 'echo x'`
  - Single command:
    - `pintos --fs-disk=2 -p ../../examples/echo -a echo -- -f -q run 'echo x'`
- `$ make check` – Builds the disk automatically
  - Copy&paste the commands make check does!

# Various directories

---

- Few user programs:
  - src/examples
- Relevant files:
  - userprog/
- Other files:
  - threads/, filesys/



# Requirements

---

- Process Termination Messages
- **Argument Passing**
- **System calls**
- Deny writes to executables

# Process Termination

---

- Process Terminates
  - `printf ("%s: exit(%d)\n", ...);`
    - Return Code
  - for eg: args-single: `exit(0)`
    - Program name
- Do not print any other message!

# Argument Passing

---

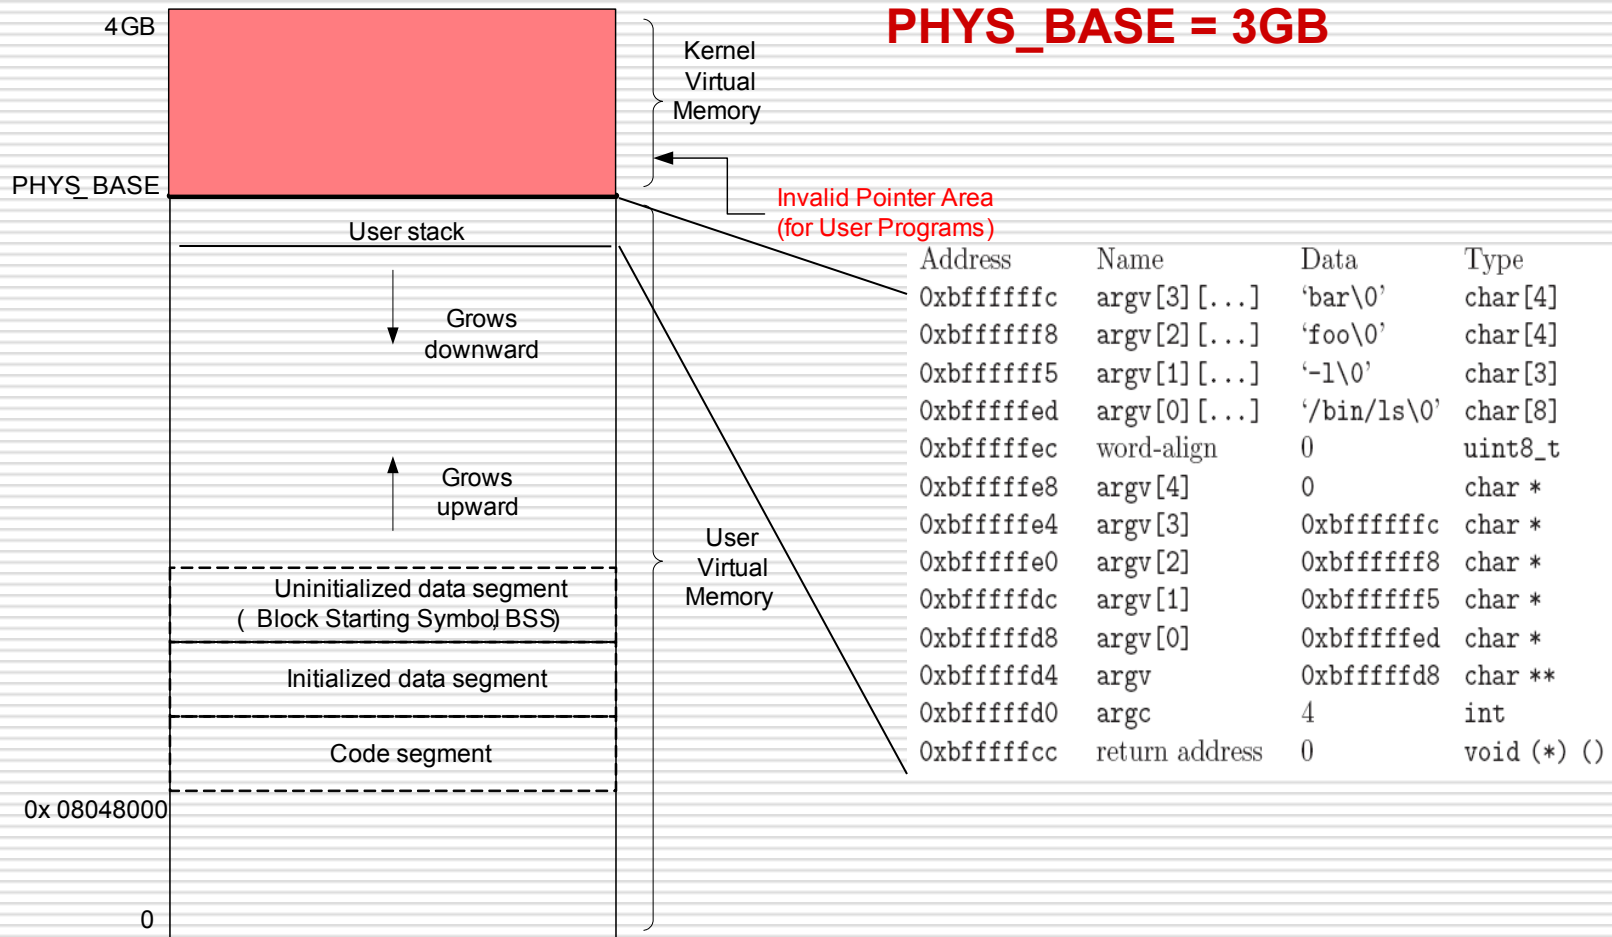
- ❑ Pintos currently lacks argument passing. You Implement it!
- ❑ Change `*esp = PHYS_BASE` to `*esp = PHYS_BASE - 12` in `setup_stack()` to get started
- ❑ Change `process_execute()` in `process.c` to process multiple arguments
- ❑ Could limit the arguments to fit in a page(4 kb)
- ❑ String Parsing: `strtok_r()` in `lib/string.h`

```
pgm.c
main(int argc,
 char *argv[]) {
 ...
}

$ pintos run 'pgm alpha beta'
argc = 3
argv[0] = "pgm"
argv[1] = "alpha"
argv[2] = "beta"
```

# Memory layout

**PHYS\_BASE = 3GB**



# Setting up the Stack

---

How to setup the stack for the program: `/bin/ls -l foo bar`

| Address      | Name           | Data         | Type        |
|--------------|----------------|--------------|-------------|
| 0xbfffffffcc | argv[3] [...]  | 'bar\0'      | char[4]     |
| 0xbfffffff8  | argv[2] [...]  | 'foo\0'      | char[4]     |
| 0xbfffffff5  | argv[1] [...]  | '-l\0'       | char[3]     |
| 0xbffffffed  | argv[0] [...]  | '/bin/ls\0'  | char[8]     |
| 0xbffffffec  | word-align     | 0            | uint8_t     |
| 0xbffffffe8  | argv[4]        | 0            | char *      |
| 0xbffffffe4  | argv[3]        | 0xbfffffffcc | char *      |
| 0xbffffffe0  | argv[2]        | 0xbfffffff8  | char *      |
| 0xbffffffdc  | argv[1]        | 0xbfffffff5  | char *      |
| 0xbffffffd8  | argv[0]        | 0xbffffffed  | char *      |
| 0xbffffffd4  | argv           | 0xbffffffd8  | char **     |
| 0xbffffffd0  | argc           | 4            | int         |
| 0xbffffffcc  | return address | 0            | void (*) () |

# Setting up the Stack... Contd


bfffffff0 00 00 00 00 | ....|

bffffffd0 04 00 00 00 d8 ff ff bf-ed ff ff bf f5 ff ff bf |.....|

bffffffe0 f8 ff ff bf fc ff ff bf-00 00 00 00 2f 62 69 |...../bil

bffffff0 6e 2f 6c 73 00 2d 6c 00-66 6f 00 62 61 72 00 |n/ls.-l.foo.bar.|

| Address     | Name           | Data        | Type        |
|-------------|----------------|-------------|-------------|
| 0xbffffffc  | argv[3][...]   | 'bar\0'     | char[4]     |
| 0xbffffff8  | argv[2][...]   | 'foo\0'     | char[4]     |
| 0xbffffff5  | argv[1][...]   | '-l\0'      | char[3]     |
| 0xbffffffd  | argv[0][...]   | '/bin/ls\0' | char[8]     |
| 0xbfffffec  | word-align     | 0           | uint8_t     |
| 0xbffffffe8 | argv[4]        | 0           | char *      |
| 0xbffffffe4 | argv[3]        | 0xbffffffc  | char *      |
| 0xbffffffe0 | argv[2]        | 0xbffffff8  | char *      |
| 0xbffffffdc | argv[1]        | 0xbffffff5  | char *      |
| 0xbffffffd8 | argv[0]        | 0xbffffffd  | char *      |
| 0xbffffffd4 | argv           | 0xbffffffd8 | char **     |
| 0xbffffffd0 | argc           | 4           | int         |
| 0xbfffffcc  | return address | 0           | void (*) () |



# Synchronization

---

- Synchronization between parent and children processes
  - Ensuring child process Loading new executables successfully

# Requirements

---

- Process Termination Messages
- **Argument Passing**
- **System calls**
- Deny writes to executables



# System Calls

---

- ❑ Pintos lacks support for system calls currently!
- ❑ Implement the system call handler in `userprog/syscall.c`
- ❑ System call numbers defined in `lib/syscall-nr.h`
- ❑ Process Control: `exit`, `exec`, `wait`
- ❑ File system: `create`, `remove`, `open`, `filesize`, `read`, `write`, `seek`, `tell`, `close`
- ❑ Others: `halt`

Syscall handler currently ...


```
static void
syscall_handler (struct intr_frame *f
UNUSED)
{
 printf ("system call!\n");
 thread_exit ();
}
```

# Continued...

- A system call has:
  - System call number
  - (possibly) arguments
- When `syscall_handler()` gets control:

```
syscall_handler (struct intr_frame *f) {
 f->esp _____

 f->eax = ... ;
}
```



Caller's User Stack

- System calls that return a value ()  
must modify  
**f->eax**

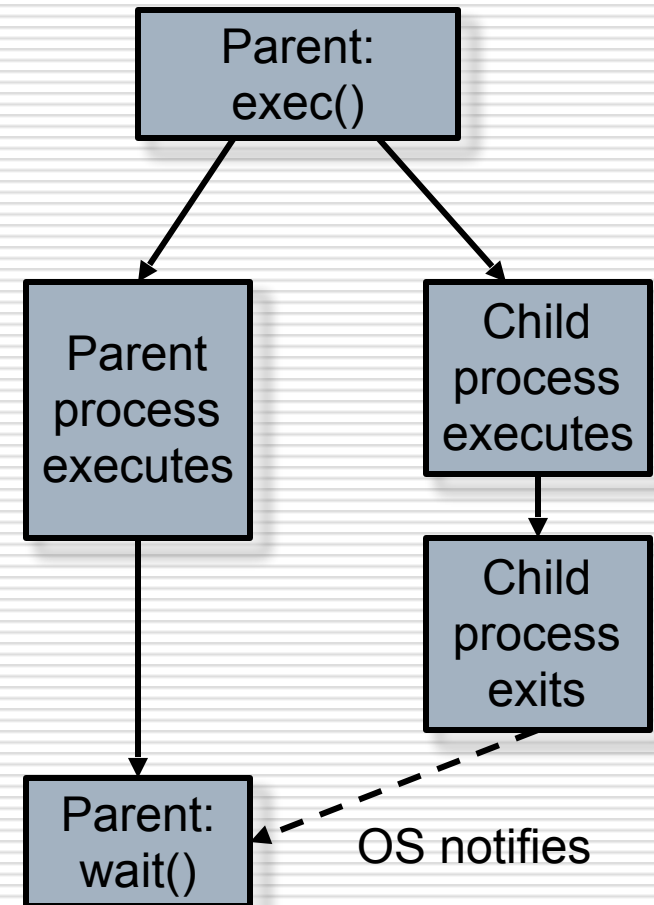
# System calls – File system

---

- Decide on how to implement the file descriptors
  - $O(n)$  data structures... perfectly fine!
- Access granularity is the entire file system
  - Have 1 global lock!
- `write()` – fd 1 writes to console
  - use `putbuf()` to write entire buffer to console
- `read()` – fd 0 reads from console
  - use `input_getc()` to get input from keyboard
- Implement the rest of the system calls

# System calls – Process Control

- `wait(pid)` – Waits for process `pid` to die and returns the status `pid` returned from `exit`
- Returns -1 if
  - `pid` was terminated by the kernel
  - `pid` does not refer to child of the calling thread
  - `wait()` has already been called for the given `pid`
- `exec(cmd)` – runs the executable whose name is given in command line
  - returns -1 if the program cannot be loaded
- `exit(status)` – terminates the current program, returns `status`
  - status of 0 indicates success, non zero otherwise



# Process Control: wait

- Implement process\_wait() in process.c
- Then, implement wait() using process\_wait()
- Cond variables and/or semaphores will help
  - Think about what semaphores may be used for and how they must be initialized
- Some Conditions to take care!
  - Parent may or may not wait for its child
  - Parent may call wait() after child terminates!

```
int
process_wait (tid_t
child_tid UNUSED)
{
 return -1;
}
```

```
main() {
int i; pid_t p;
p = exec("pgm a b");
// i = wait (p);
}
```

# Memory Access

---

- System calls can have memory access
  - e.g open(), read(), write()
  - have a look at tests cases \*-bad-ptr.c
- In open-bad-ptr.c
  - ```
void test_main (void)
{
    msg ("open(0x20101234): %d",
        open ((char *) 0x20101234));
    fail ("should have called exit(-1)");
}
```

Memory Access (contd')

- ❑ Invalid pointers must be rejected. Why?
 - Kernel has access to all of physical memory including that of other processes
 - Kernel like user process would fault when it tries to access unmapped addresses
- ❑ User process cannot access kernel virtual memory
- ❑ User Process after it has entered the kernel can access kernel virtual memory and user virtual memory
- ❑ How to handle invalid memory access?

Memory Access (contd')

- Two methods to handle invalid memory access
 - Verify the validity of user provided pointer and then dereference it
 - Look at functions in `userprog/pagedir.c`, `threads/vaddr.h`
 - Strongly recommended!
 - Check if user pointer is below `PHYS_BASE` and dereference it
 - Could cause page fault
 - Handle the page fault by modifying the `page_fault()` code in `userprog/exception.c`
 - Make sure that resources are not leaked

Some Issues to look at...

- ❑ Check the validity of the system call parameters
- ❑ Every single location should be checked for validity before accessing it. For e.g. not only `f->esp`, but also `f->esp + 1`, `f->esp+2` and `f->esp+3` should be checked
- ❑ Read system call parameters into kernel memory (except for long buffers)
 - `copy_in` function recommended!

Denying writes to Executables

- ❑ Use `file_deny_write()` to prevent writes to an open file
- ❑ Use `file_allow_write()` to re enable write
- ❑ Closing a file will automatically re enable writes

Suggested Order of Implementation

- ❑ Change `*esp = PHYS_BASE` to `*esp = PHYS_BASE - 12` to get started
- ❑ Implement the system call infrastructure
- ❑ Change `process_wait()` to a infinite loop to prevent pintos getting powered off before the process gets executed
- ❑ Implement `exit` system call
- ❑ Implement `write` system call
- ❑ Start making other changes

Misc

- ❑ Deadline: 19 Oct, 11:59 pm
- ❑ Do not forget the design document
 - Must be done individually
- ❑ Good Luck!